

Lab 6 - Cache

CSED 311 Computer Architecture Lab

Sungjun Cho

2021 - 05 - 11

allencho1222@postech.ac.kr

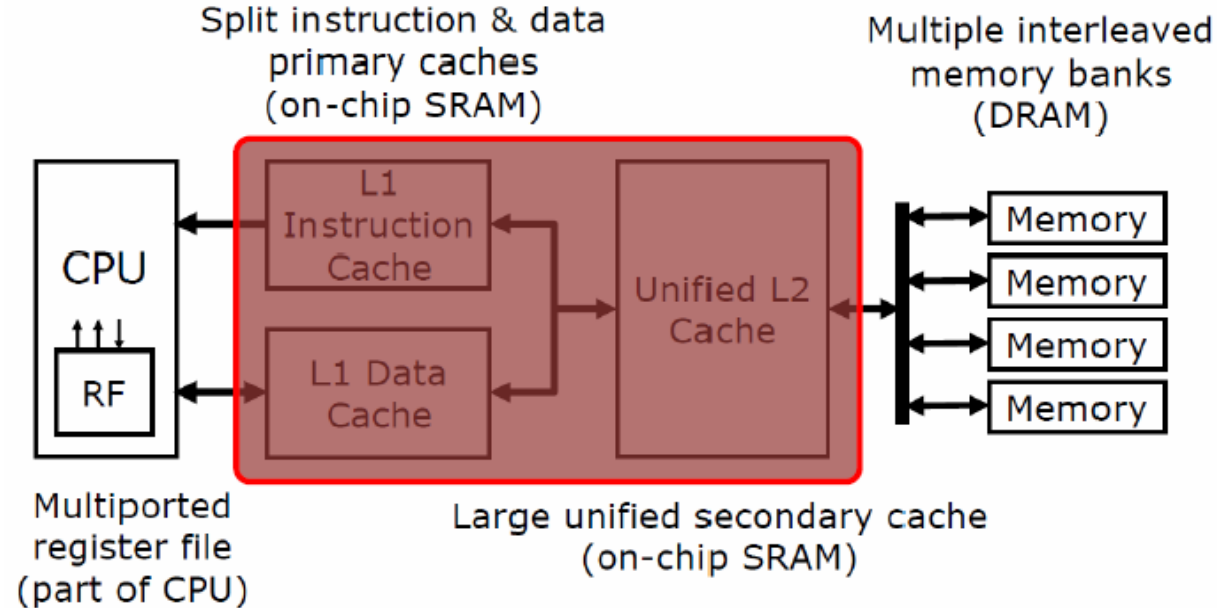
csed311-ta@postech.ac.kr

Objectives

- Understand how cache works
- Implement a set-associative cache on your pipelined CPU
- Evaluate the speed-up achieved by using cache
 - Hit ratio
 - Corresponding speed-up (vs. no-cache CPU)

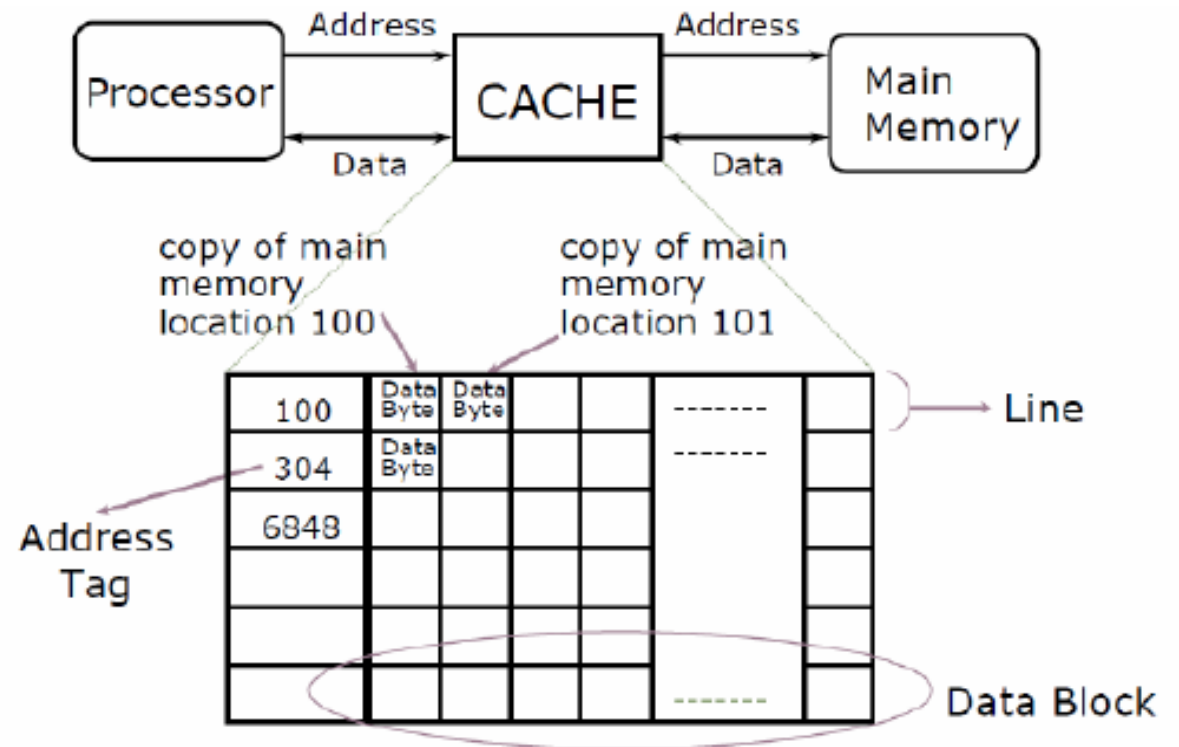
Cache

- Mitigating the gap between CPU and memory
 - Memory access: few hundred cycles
 - Cache access: few cycles
- Why does it work?
 - Locality!





Cache: internal structure

- Tag
 - Detect address conflicts
- Data
 - Fetch by line: exploiting spatial locality



Cache: associativity

- Associativity
 - Reducing address conflicts
- Direct mapped, n-way, fully associative
 - Tradeoffs exist



Cache: other design choices

- Replacement policy
 - Random, LRU, FIFO, ...,
 - Each one has strengths & drawbacks
- Write policy
 - Write-through, writeback, write-no-allocate
 - Related to coherency management

Lab Assignment 06 (1/5)

- Design and implement your own cache with the following requirements:
 - 2-way set associative, single-level cache
 - Capacity: 32 words / Line size: 4 words
 - If hit, return data in the following cycle
 - It should be a part of CPU (not TB)
- You have to choose the following design choices:
 - Replacement policy, write policy
 - Unified (32 words == 8 Lines) or separate I/D (4 lines ea.) cache

Lab Assignment 06 (2/5)

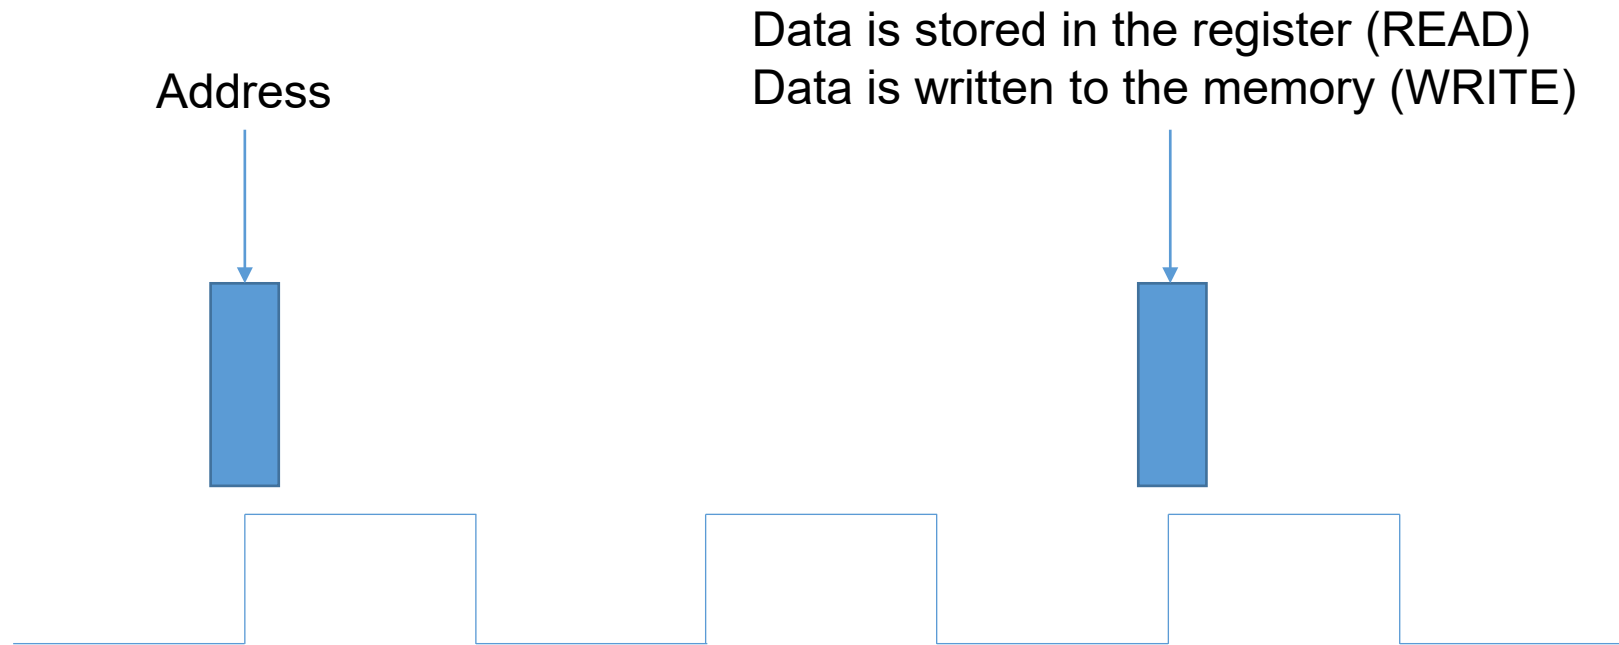
- You may implement direct mapped cache with reduced score of 80%
 - Direct-mapped, single-level cache
 - Capacity: 16 words / Line size: 4 words
 - If hit, return data in the following cycle
 - It should be a part of CPU (not TB)

Lab Assignment 06 (3/5)

- Memory access latency – Previous pipelined CPU
 - One memory access fetches one word with one cycle
- New memory access latency
 - You're required to modify memory.v
 - Cache hit takes **one cycle**
 - One memory access should fetch **four words into cache (= one cache line) and take six cycles**
 - For your baseline CPU (CPU without cache), one memory access should fetch **one word** and take **two cycles**
 - Then, you need to implement two different memory models
 - ✓ Memory for CPU with cache: return four words in six cycles
 - ✓ Memory for CPU without cache: return one word in two cycles


Waveform (Baseline CPU)

- Cache does not exist




Waveform (Cache hit)

- Cache miss/hit and data can be checked in a single cycle

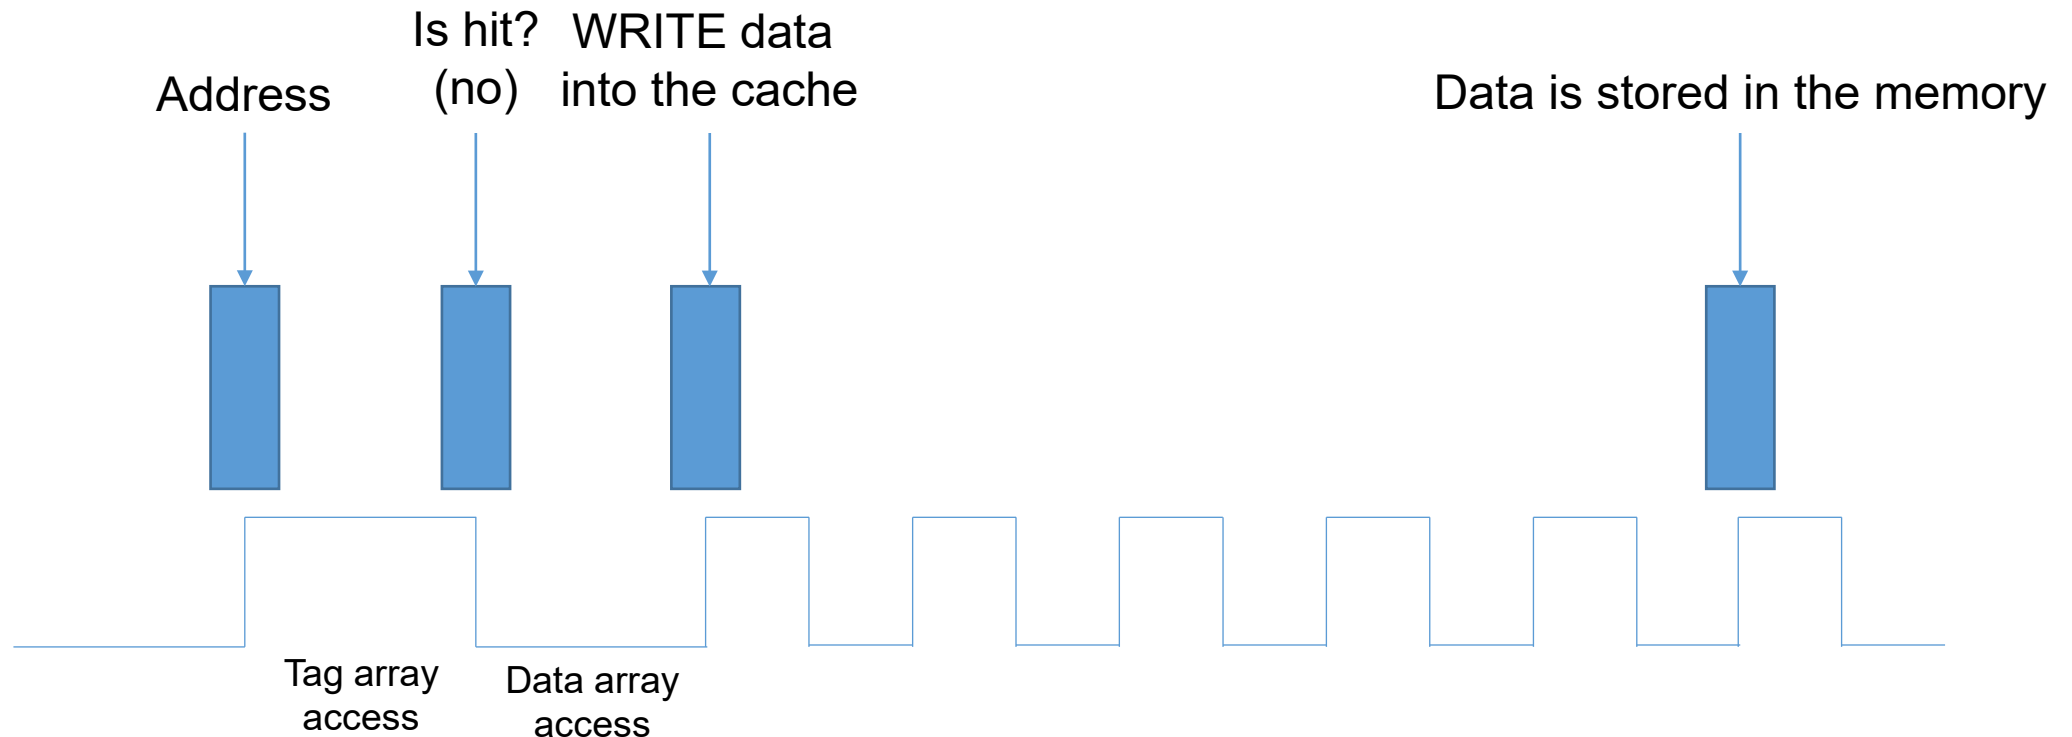


Waveform (Cache miss)




Waveform (WRITE-THROUGH)

- Usually comes with write-no-allocate
 - **Cache hit** (WRITE request)
 - ✓ WRITE data into the cache + WRITE data into the memory



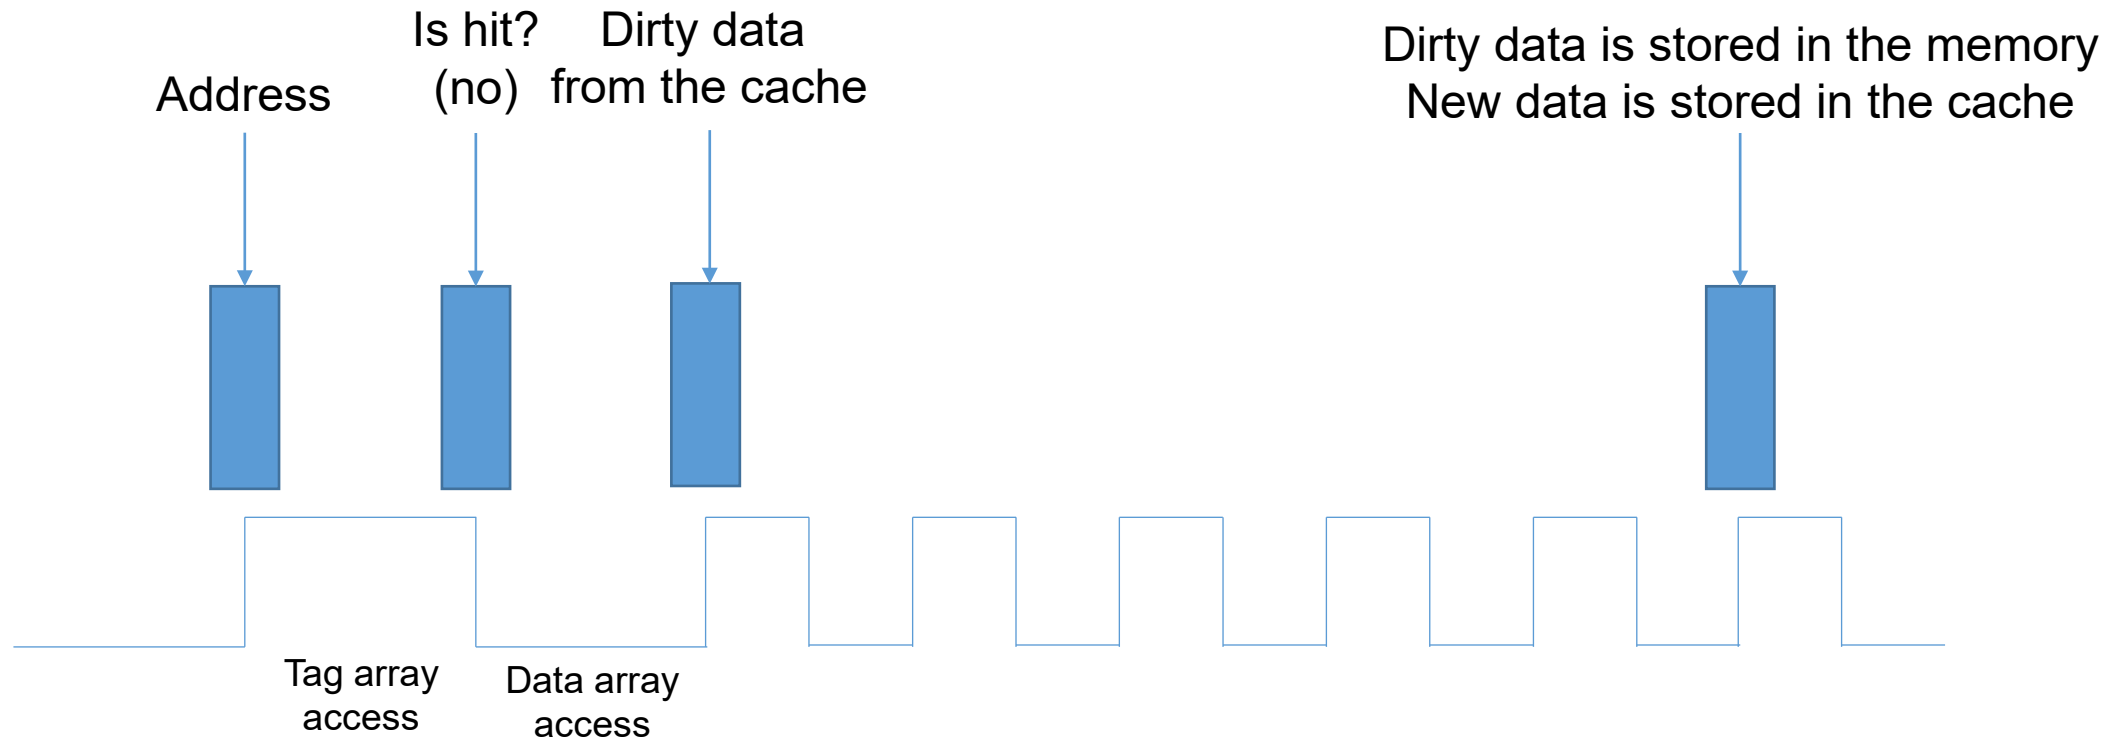
Waveform (WRITE-THROUGH)

- Usually comes with write-no-allocate
 - **Cache miss** (WRITE request)
 - ✓ WRITE data into the memory



Waveform (WRITE-BACK)

- Usually comes with write-allocate
 - **Cache miss** (WRITE request)
 - ✓ READ data from the memory + WRITE dirty data into the memory



Lab Assignment 06 (4/5)

- Memory requirements
 - You can use either a 2-port RAM or a single-port RAM
 - Latencies of RAM
 - ✓ Should be serialized



⇒ Different ports can handle independent requests!

Lab Assignment 06 (5/5)

- For the report, the following contents should be included
 - Describe your design choice
 - Calculate the hit (or miss) ratio
 - ✓ Hit ratio = $(\# \text{ of hits}) / (\# \text{ of memory accesses})$
 - Compare the performance
 - ✓ CPU without cache vs. CPU with cache
 - ✓ You should use the new latencies!

Announcement

- A skeleton code from this lab is not provided.
You should implement cache *on your previous implementation of pipelined CPU*.
- You *cannot begin* this assignment unless you have finished previous work.